

Objective

- We implemented a Graph Neural Network (GNN)-based Movie Recommendation System using the MovieLens dataset.
- This system models movies as nodes and user interactions (ratings) as edges in a graph.
- The goal was to predict movie ratings based on user behavior using a deep learning model while optimizing for speed and efficiency.

What We Did & Why? (Real-World Application Perspective)

1 Constructed a Graph from User-Movie Interactions

What we did: Converted the MovieLens dataset into a graph, where: Movies = Nodes User interactions (ratings) = Edges (connecting movies rated by the same user) Why?

This allows us to capture relationships between movies (e.g., if two movies are frequently watched by the same users, they are linked). This is similar to how Netflix or YouTube group content based on shared audience preferences.

2 Built a GNN Model for Rating Prediction

What we did: Implemented a Graph Convolutional Network (GCN) to learn movie relationships.

Why?

Traditional recommendation systems (collaborative filtering) fail when a new movie has no ratings (cold start problem).

GNNs generalize better by learning from graph structure instead of relying only on raw ratings.

This approach is similar to how LinkedIn suggests new connections based on mutual connections.

3 Optimized Model for Faster Execution

What we did:

Reduced dataset size to speed up processing.

Used adjacency matrices instead of loops for efficient graph construction.

Simplified the model architecture (GCN instead of GAT) to reduce computation.

Implemented early stopping to prevent unnecessary training.

Why?

Real-world relevance: In platforms like Netflix, Amazon, and Spotify,

recommendation models need to process millions of users quickly. Speed optimizations ensure faster recommendations without sacrificing accuracy.

```
In [2]: import torch
import torch.nn.functional as F
from torch_geometric.data import Data
from torch_geometric.nn import GCNConv
import pandas as pd
import numpy as np
import networkx as nx
from sklearn.model_selection import train_test_split
from torch_geometric.utils import from_networkx

# Load Data (Limit for Faster Processing)
movies = pd.read_csv('movies.csv').head(5000) # Limit dataset for spe
ratings = pd.read_csv('ratings.csv').head(10000)

# Create Movie ID Mappings
movie_id_map = {id_: idx for idx, id_ in enumerate(movies['movieId'].u
num_movies = len(movie_id_map)

# Create Graph using NetworkX
G = nx.Graph()

# Add Movie Nodes
G.add_nodes_from(movie_id_map.values())

# Add Edges (Efficient Adjacency Matrix Representation)
user_movie_map = ratings.groupby('userId')['movieId'].apply(list)

for movies_watched in user_movie_map:
    movie_indices = [movie_id_map[m] for m in movies_watched if m in m
    G.add_edges_from([(movie_indices[i], movie_indices[j]) for i in ra

# Convert Graph to PyTorch Geometric Format
data = from_networkx(G)

# Assign Features (Smaller Feature Dimension for Speed)
data.x = torch.rand((num_movies, 8)) # Reduced to 8 dimensions

# Assign Labels (Movie Ratings as a Regression Task)
avg_ratings = ratings.groupby('movieId')['rating'].mean()
labels = np.array([avg_ratings.get(movies.iloc[i]['movieId'], 2.5) for
data.y = torch.tensor(labels, dtype=torch.float32).view(-1, 1)

# Train/Test Split (Efficient Subsampling)
train_idx, test_idx = train_test_split(np.arange(num_movies), test_siz
data.train_mask = torch.tensor(np.isin(np.arange(num_movies), train_id
data.test_mask = torch.tensor(np.isin(np.arange(num_movies), test_idx)
```

```

# Define Optimized GNN Model (Smaller Architecture)
class GNN(torch.nn.Module):
    def __init__(self, in_channels, hidden_channels, out_channels):
        super(GNN, self).__init__()
        self.conv1 = GCNConv(in_channels, hidden_channels)
        self.conv2 = GCNConv(hidden_channels, out_channels) # Remove

    def forward(self, x, edge_index):
        x = self.conv1(x, edge_index).relu()
        x = self.conv2(x, edge_index)
        return x

# Initialize Model, Optimizer, and Device
device = torch.device('cuda' if torch.cuda.is_available() else 'cpu')
model = GNN(in_channels=8, hidden_channels=16, out_channels=1).to(device)
optimizer = torch.optim.Adam(model.parameters(), lr=0.005)

data = data.to(device)

# Train with Early Stopping
best_loss = float('inf')
patience, patience_counter = 5, 0 # Early stopping params

for epoch in range(50): # Reduced epochs
    model.train()
    optimizer.zero_grad()
    out = model(data.x, data.edge_index)
    loss = F.mse_loss(out[data.train_mask], data.y[data.train_mask])
    loss.backward()
    optimizer.step()

    # Early Stopping Check
    if loss.item() < best_loss:
        best_loss = loss.item()
        patience_counter = 0
    else:
        patience_counter += 1
        if patience_counter >= patience:
            print(f"Stopping early at epoch {epoch}")
            break

    if epoch % 10 == 0:
        print(f"Epoch {epoch}, Loss: {loss.item():.4f}")

# Evaluate Model
model.eval()
with torch.no_grad():
    test_loss = F.mse_loss(model(data.x, data.edge_index)[data.test_mask], data.y[data.test_mask])
    print(f"Test MSE Loss: {test_loss.item():.4f}")

```

Epoch 0, Loss: 8.4611
Epoch 10, Loss: 4.7286
Epoch 20, Loss: 2.2814
Epoch 30, Loss: 1.0573
Stopping early at epoch 39
Test MSE Loss: 1.0960